Crimea Turns Out to Have Been Ours

Measuring Authoritarian Narratives

Perry Carter

June 26, 2025

New York University, Abu Dhabi

Outline

Motivation

- Narratives are at the core of many theories of politics.
- LLM-fueled explosion in analysis of text
- Topic modeling provides a flexible, easy-to-use method to study content.
- Yet, study of narratives still relies on categorical labels.

Primary Aim
Provide an interpretable, principled method for representing narratives in both elite and mass speech.

Validation: A Propaganda Shock in Russian Media

- March 2014: news aggregator website Lenta.ru publishes an interview with a Ukrainian nationalist
- Within days, warning from the state censor $\rightarrow$ mass editorial purge (Golitsyna, 2014)
- Result was rapid alignment with Kremlin perspectives on protest, legitimacy, geopolitics
    - Still writes (mostly) about the same things – just offers a different causal worldview.

Repression Event and Daily Protest Coverage

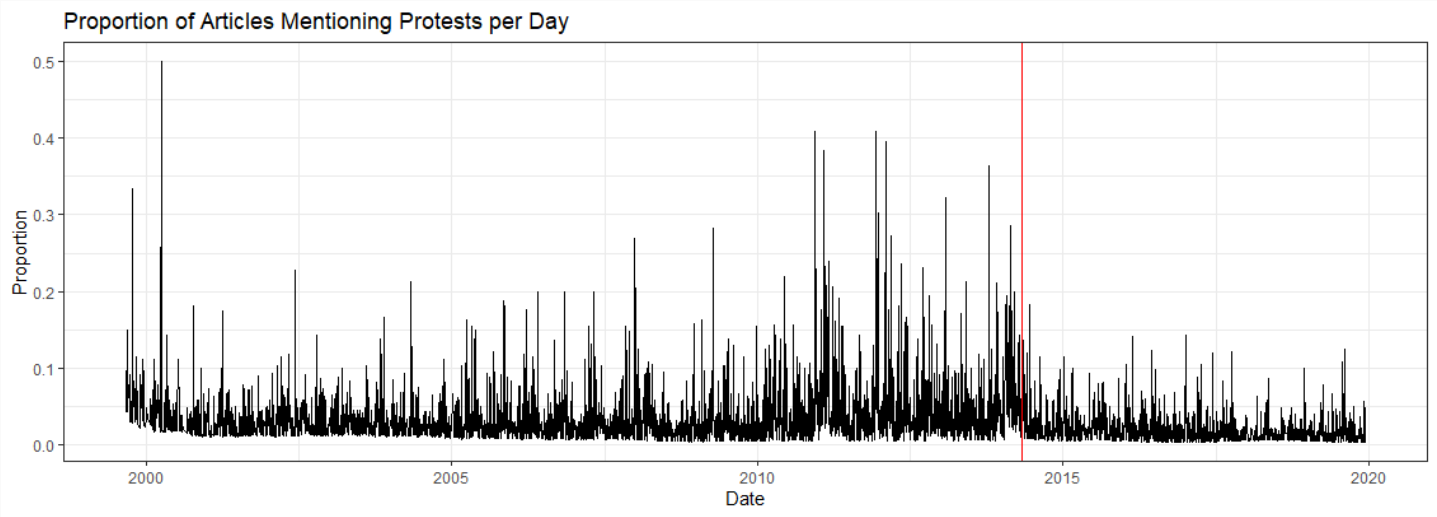

Figure 1: Daily share of protest-related articles on Lenta.ru

Censorship Effects by Topic

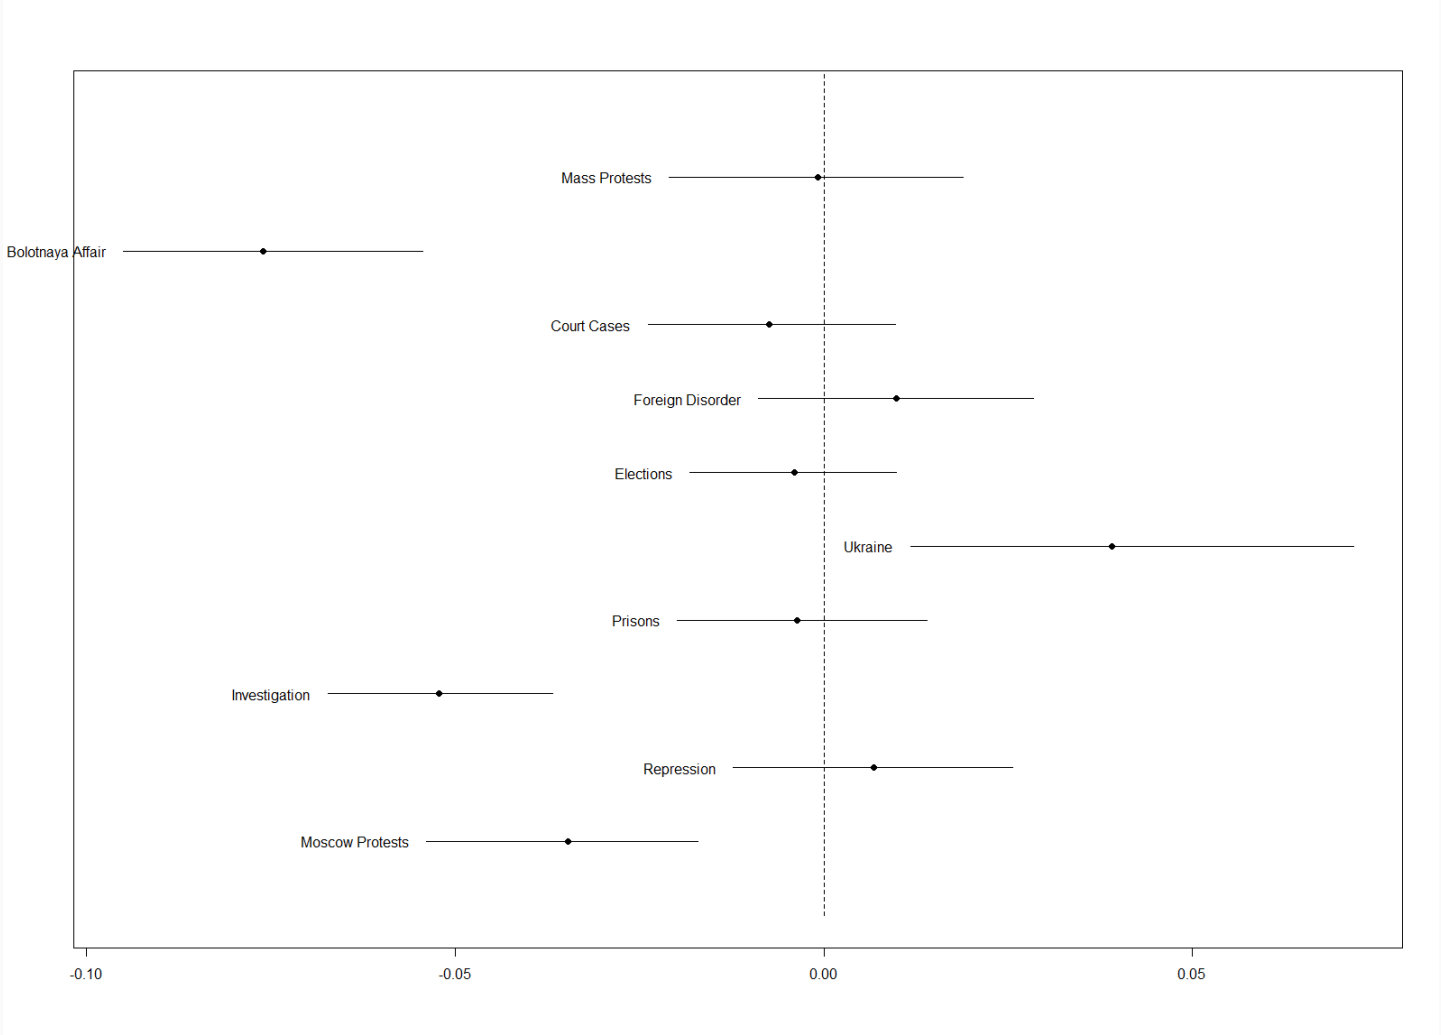

Figure 2: RDiT estimates for major topic categories

From Text to Narrative Graphs

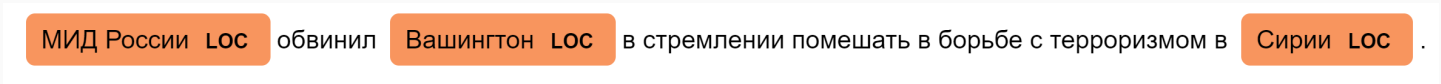

Figure 3: Example Sentence Representation [The Ministry of Foreign Affairs accused Washington of trying to interfere with the fight against terrorism in Syria]

Extracting Latent Roles

- We treat each entity (e.g. "Navalny," "demonstrators," "sign") as a node in a weighted directed graph (per document).
- Core intuition: each document constitutes an instance of a meta-narrative in which entities take on narrative roles (e.g. villains).
- Define a signed stochastic block model groups nodes into K *roles* so that
  - Nodes in the same role have similar patterns of outgoing and incoming edges
  - Roles capture narrative functions (e.g. leaders, actions, contexts)
- The model estimates (through semi-collapsed variational EM)
  1. A **role assignment** for each entity
  2. A **connection probability matrix** $B_{k\ell}$ governing how often role k speaks about role ℓ
  3. A **polarity matrix** $\mu_{k\ell}$ gives positive/negative framing between roles
  4. Document and dyadic covariate influence estimates with exact credible intervals.

Interpreting Role Interactions and Change

- The fitted B matrix tells us, for each ordered pair of roles, how strongly they interact
  - e.g. a high $B_{\text{Leaders, ProtestActors}}$ means leaders frequently "act on" protestors
- Polarity parameters show whether those interactions are framed positively or negatively
- By comparing pre- vs. post-purge estimates of B and polarity, we map how narrative structure shifts
- Example takeaway: after censorship, the "Leaders" role shifts from directing action to primarily *responding* to protestors

Results: Latent Roles in Protest Narratives

- **0: Leaders** (e.g. Navalny, Putin) — "Navalny accused Putin of..."
- **1: Events/Actions** (declare, sign, justice) — "The parties declared their intention to..."
- **2: Context/Foreigners** (London, Armenia, election) — e.g. foreign statements contextualizing protests
- **3: Protest Actors** (demonstrators, Muslims) — "Demonstrators in London demanded freedom and justice..."

Effects of Censorship on Role Interactions

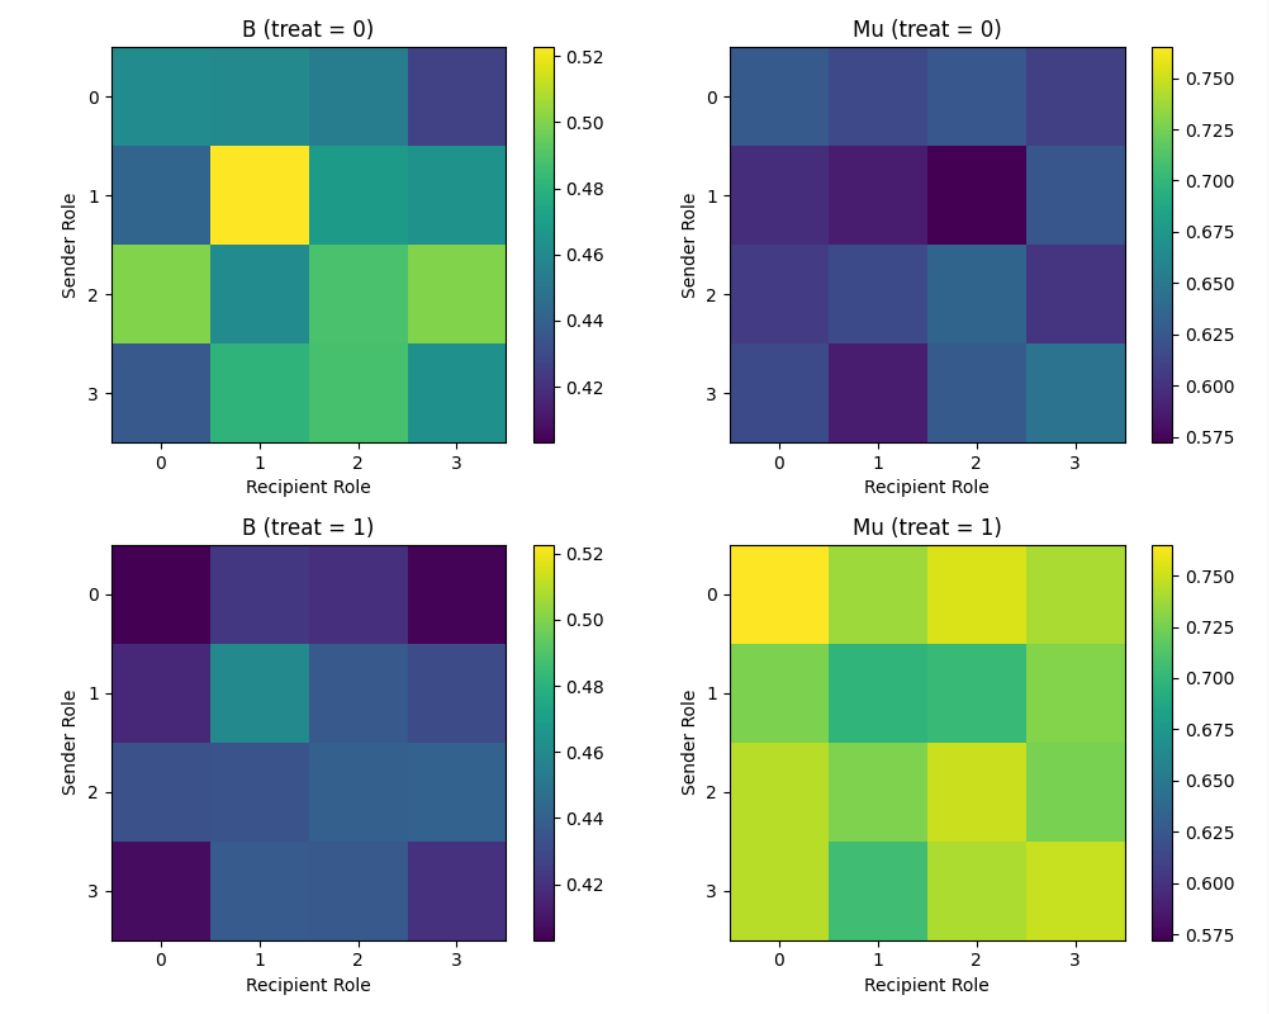

Figure 4: Estimated change in interaction probabilities and positivity after the purge (treat = 1).

Contributions & Next Steps

- I introduce a scalable, interpretable, and transparent narrative-graph framework for estimating changes in narratives.
- A key advantage is domain independence: mass and elite representation admit similar styles after controlling for baseline entitivity.
- **Future work**: Apply to larger corpora and incorporate dyadic context (setting, tense).

Thank you!

Please send comments and questions by email to pjc504@nyu.edu